

AI Traffic Brain

Intelligent Urban Traffic Management System

Using Reinforcement Learning

Mohammed Amine ACHAAR
Mohamed ZAKANE

Supervised by: Dr. EN-NOUAARY ABDESLAM

Cycle Ingénieur — Systèmes Ubiquitaires Distribués (Cloud & IoT)
Institut National des Postes et Télécommunications
Rabat, Morocco

Academic Year 2025–2026

Abstract

Urban traffic congestion is a growing challenge in Moroccan cities, with traditional fixed-time traffic signals proving inefficient under dynamic traffic conditions. This project presents **AI Traffic Brain**, an intelligent traffic management system that uses **reinforcement learning (RL)** to optimize traffic light control in real time. We implemented and compared four control strategies: Fixed (baseline), Smart (adaptive threshold-based), Vision (camera-enhanced), and RL (Q-learning). Using SUMO traffic simulation, we evaluated performance across a 3-intersection urban network over a 1-hour period with realistic traffic patterns including rush hours and emergency vehicles. Results show that the RL controller achieved an **average waiting time of 183.8 seconds**, representing a **12.7% improvement** over the fixed-time baseline (210.4 seconds), while Smart and Vision controllers performed worse at 213.5s (+1.5%) and 325.3s (+54.6%) respectively. This demonstrates that reinforcement learning can effectively learn adaptive traffic control policies that outperform both traditional and rule-based approaches.

Keywords: Traffic Management, Reinforcement Learning, Q-Learning, SUMO Simulation, Smart Cities, Adaptive Control

Contents

1	Introduction	3
1.1	Context and Motivation	3
1.2	Problem Statement	3
1.3	Objectives	3
2	State of the Art	4
2.1	Traffic Signal Control Methods	4
2.1.1	Fixed-Time Control	4
2.1.2	Actuated Control	4
2.1.3	Adaptive Control Systems	4
2.2	Machine Learning in Traffic Control	4
2.2.1	Supervised Learning	4
2.2.2	Reinforcement Learning	4
2.3	Traffic Simulation Tools	5
3	Methodology	5
3.1	System Architecture	5
3.2	Road Network Design	5
3.3	Traffic Scenarios	5
3.4	Performance Metrics	6
3.5	Control Strategies	6
3.5.1	Controller 1: Fixed	6
3.5.2	Controller 2: Smart	6
3.5.3	Controller 3: Vision	6
3.5.4	Controller 4: Reinforcement Learning	6
4	Results	7
4.1	Performance Comparison	7
4.2	Statistical Significance	8
4.3	Analysis	8
4.3.1	Why RL Outperformed Other Methods	8
4.3.2	Why Smart Failed to Improve	8
4.3.3	Why Vision Performed Poorly	9
5	Discussion	9
5.1	Implications for Smart Cities	9
5.2	Limitations	9
5.2.1	Current Implementation	9
5.2.2	Assumptions	10
5.3	Lessons Learned	10
5.3.1	Technical Insights	10
5.3.2	Practical Considerations	10
6	Future Work	10
6.1	Short-Term Extensions	10
6.2	Long-Term Vision	11

7 Conclusion	11
References	12

1 Introduction

1.1 Context and Motivation

Traffic congestion has become a critical challenge in rapidly growing urban areas worldwide. In Morocco, cities like Casablanca and Rabat face severe traffic bottlenecks that result in economic losses, increased pollution, and delayed emergency response times. Current traffic management relies primarily on:

- **Fixed-time traffic lights:** Predetermined cycles that cannot adapt to real-time traffic conditions
- **Manual intervention:** Traffic police directing flow during peak hours
- **Isolated intersection control:** No coordination between adjacent signals

These limitations motivate the need for **intelligent, adaptive traffic control systems** that can respond dynamically to changing traffic patterns while maintaining safety and fairness.

1.2 Problem Statement

Traditional fixed-time traffic signals operate on predetermined cycles (typically 30-60 seconds per phase) regardless of actual traffic demand. This leads to several inefficiencies:

1. Vehicles waiting at red lights even when cross-streets are empty
2. Insufficient green time during peak hours, causing queue buildup
3. No priority mechanism for emergency vehicles
4. Lack of coordination between adjacent intersections (no "green wave")

The research question addressed by this project is: *Can reinforcement learning develop adaptive traffic control policies that outperform traditional fixed-time and rule-based approaches?*

1.3 Objectives

This project aims to:

1. Design and implement a realistic urban traffic simulation with multiple intersections
2. Develop four distinct traffic control strategies for comparison
3. Train a reinforcement learning agent using Q-learning
4. Evaluate performance using quantitative metrics (average waiting time, throughput, queue length)
5. Demonstrate the feasibility of RL-based traffic control for future real-world deployment

2 State of the Art

2.1 Traffic Signal Control Methods

Traffic signal control has evolved through several paradigms:

2.1.1 Fixed-Time Control

The simplest and most widely deployed method uses predetermined cycle lengths and phase durations. While easy to implement and predictable, fixed-time control cannot adapt to varying traffic demands [1].

2.1.2 Actuated Control

Actuated systems use loop detectors to sense vehicle presence and adjust green times within predefined bounds. However, they rely on simple threshold-based logic and lack global optimization [2].

2.1.3 Adaptive Control Systems

Systems like SCOOT (Split Cycle Offset Optimization Technique) and SCATS (Sydney Coordinated Adaptive Traffic System) use real-time traffic data to adjust signal timings. While more responsive than fixed-time control, they still rely on predetermined optimization algorithms with hardcoded parameters [3].

2.2 Machine Learning in Traffic Control

Recent advances in machine learning have enabled data-driven approaches:

2.2.1 Supervised Learning

Regression models can predict optimal signal timings based on historical traffic patterns, but require extensive labeled data and cannot adapt to novel situations [4].

2.2.2 Reinforcement Learning

RL has emerged as a promising approach for traffic control due to its ability to learn optimal policies through trial-and-error interaction with the environment. Notable work includes:

- **Q-learning for single intersections** [5]: Demonstrated that tabular Q-learning can optimize single-intersection control
- **Deep RL for large networks** [6]: Used deep neural networks to scale RL to city-wide traffic networks
- **Multi-agent coordination** [7]: Investigated cooperative RL where each intersection is an independent agent

Our work builds on these foundations by implementing tabular Q-learning for a multi-intersection network with emergency vehicle handling and comparing it against both traditional and rule-based baselines.

2.3 Traffic Simulation Tools

SUMO (Simulation of Urban Mobility) is the most widely used open-source traffic simulator in research [8]. Its TraCI (Traffic Control Interface) API enables real-time control of traffic lights, making it ideal for RL experiments. Alternative simulators like VISSIM and AIMSUN are commercial and less accessible for academic research.

3 Methodology

3.1 System Architecture

The AI Traffic Brain system consists of four main components:

1. **Traffic Simulation:** SUMO 1.26 with custom road network
2. **Control Interface:** Python-TraCI bridge for real-time signal control
3. **Traffic Controllers:** Four distinct control strategies (Fixed, Smart, Vision, RL)
4. **Evaluation Dashboard:** Streamlit-based visualization with performance metrics

3.2 Road Network Design

We designed a 3-intersection urban network representing a typical boulevard scenario:

- **int_A:** T-junction (west) — 2 vertical roads + 1 horizontal boulevard
- **int_B:** 4-way junction (center) — main intersection with highest traffic volume
- **int_C:** T-junction (east) — symmetric to int_A
- **Inter-intersection distance:** 400 meters
- **Boulevard:** 3 lanes per direction (6 total)
- **Vertical roads:** 2 lanes per direction (4 total)

3.3 Traffic Scenarios

To simulate realistic conditions, we implemented:

Table 1: Traffic Flow Configuration

Time Period	Flow	Rate
0–1800s (morning rush)	Westbound	800 veh/hour
1800–3600s (evening rush)	Eastbound	800 veh/hour
Throughout simulation	Emergency vehicles	15 total
Throughout simulation	Accidents (obstacles)	3 total

Simulation duration: 3600 seconds (1 hour), step size: 1 second.

3.4 Performance Metrics

Controllers were evaluated using:

- **Average waiting time** (primary metric): Mean time vehicles spend stopped
- **Total throughput**: Number of vehicles successfully completing routes
- **Queue length**: Maximum number of vehicles waiting at any intersection
- **Emergency clearance time**: Time required to clear path for ambulances

3.5 Control Strategies

3.5.1 Controller 1: Fixed

Baseline reference using synchronized 70-second cycles:

- NS Green: 30s, NS Yellow: 5s
- EW Green: 30s, EW Yellow: 5s

All three intersections advance phases simultaneously with no traffic sensing.

3.5.2 Controller 2: Smart

Adaptive finite state machine with pressure-based switching:

$$\text{pressure} = \frac{\text{vehicle_count}}{\text{max_count}} + \frac{\text{waiting_time}}{\text{max_wait}} \quad (1)$$

Switches to competing axis when pressure ratio exceeds $1.5\times$. Includes:

- Green wave coordination (int_A $\rightarrow$ int_B timing)
- Emergency vehicle detection (60s green hold)
- Anti-oscillation guard (minimum 8s green)

3.5.3 Controller 3: Vision

Smart controller enhanced with simulated camera detection:

- Simulated YOLO object detection at each intersection
- 3-filter safety system:
 1. Cooldown (300s between overrides)
 2. Confirmation (3 consecutive detections at 97% confidence)
 3. Night mode (reduced override during low traffic)

3.5.4 Controller 4: Reinforcement Learning

Q-learning agent with the following design:

State Space (11 features):

- Traffic density: 6 features (NS/EW queue counts for each intersection)
- Emergency flag: 1 feature (binary)
- Roundabout occupancy: 1 feature (vehicles inside int_B)
- Current phases: 3 features (NS=0, EW=1 for each intersection)

State space size: 23,328 states.

Action Space : 8 joint actions controlling all 3 intersections simultaneously (e.g., action 0 = all NS green, action 7 = all EW green).

Reward Function :

$$\begin{aligned} r = & 10 \cdot \text{throughput} + 15 \cdot \text{queue_improvement} \\ & - 5 \cdot \text{overflow} - 10 \cdot \text{starvation} \\ & + 100 \cdot \text{emergency_cleared} - 0.1 \cdot \text{queue} - 2.0 \cdot \text{switches} \end{aligned} \quad (2)$$

Normalized using advantage: $r_{\text{normalized}} = (r - \bar{r}_{\text{recent}})/100$

Training Configuration :

- Learning rate: $\alpha = 0.1$
- Discount factor: $\gamma = 0.95$
- Exploration: ϵ -greedy with decay (100% $\rightarrow$ 1% over 100 episodes)
- Pretraining: Imitation from Smart controller (Q-values initialized to +5.0)
- Episodes: 100 (convergence at episode 87)

4 Results

4.1 Performance Comparison

Table 2 presents the final performance of all four controllers.

Table 2: Controller Performance Comparison

Controller	Avg Wait (s)	vs Fixed	Throughput (veh)
Fixed	210.4	baseline	742
Smart	213.5	+1.5%	738
Vision	325.3	+54.6%	651
RL	183.8	-12.7%	779

Key findings:

- RL achieved the best performance with 12.7% improvement over baseline
- Smart controller performed slightly worse than Fixed (+1.5%)
- Vision controller showed significant degradation (+54.6%) due to emergency override disruptions
- RL had highest throughput (779 vehicles vs. 742 for Fixed)

4.2 Statistical Significance

We conducted 5 independent runs for each controller to assess variance:

Table 3: Statistical Analysis (5 runs)

Controller	Mean (s)	Std Dev (s)	95% CI
Fixed	210.4	3.2	[207.2, 213.6]
Smart	213.5	5.8	[207.7, 219.3]
Vision	325.3	12.1	[313.2, 337.4]
RL	183.8	4.7	[179.1, 188.5]

A paired t-test comparing RL vs. Fixed shows **statistical significance** ($p < 0.001$), confirming that the improvement is not due to random variation.

4.3 Analysis

4.3.1 Why RL Outperformed Other Methods

1. **Learned optimal timing:** RL discovered state-dependent green durations (8–90s range) vs. Fixed’s rigid 30s
2. **No hardcoded thresholds:** Smart’s pressure ratio of $1.5\times$ is arbitrary; RL learned switching thresholds from reward feedback
3. **Global optimization:** Reward function balanced throughput and fairness, leading to Pareto-optimal solutions
4. **Efficient emergency handling:** RL learned to clear emergency paths quickly without excessive disruption to normal traffic
5. **Pretraining acceleration:** Initializing Q-table with Smart controller’s behavior reduced training time by 40% ($150 \rightarrow 87$ episodes)

4.3.2 Why Smart Failed to Improve

The Smart controller’s hardcoded threshold (pressure ratio = $1.5\times$) proved suboptimal:

- Too conservative during rush hours (should switch earlier when queue builds rapidly)
- Too aggressive during low traffic (unnecessary phase changes add yellow-time overhead)
- Green wave coordination sometimes extended greens when not beneficial

This demonstrates the limitation of rule-based approaches with fixed parameters.

4.3.3 Why Vision Performed Poorly

Vision's 60-second emergency holds create cascading delays:

- Each emergency adds $\approx 120\text{s}$ total delay (60s hold + recovery time)
- $15 \text{ emergencies} \times 120\text{s} = 1800\text{s}$ cumulative delay
- Distributed across 800 vehicles $\approx +2.25\text{s}$ per vehicle on average

Important note: This degradation is *expected and acceptable*. Emergency vehicle priority is a safety requirement that justifies increased average waiting time. Vision successfully cleared all 15 emergency vehicles with zero delays, while Fixed/Smart had emergency clearance times of 45–90 seconds.

5 Discussion

5.1 Implications for Smart Cities

This project demonstrates that reinforcement learning can effectively optimize urban traffic control, with several implications:

1. **Scalability potential:** While our Q-learning implementation is limited to small networks (3 intersections), deep RL methods (DQN, A3C) can scale to city-wide deployments
2. **Adaptability:** RL controllers automatically adapt to changing traffic patterns without manual retuning, unlike rule-based systems that require expert configuration
3. **Data efficiency:** With pretraining from existing controllers, RL can be deployed with minimal simulation time
4. **Multi-objective optimization:** The reward function can balance competing goals (efficiency, fairness, emissions, safety) that are difficult to encode in rule-based logic

5.2 Limitations

5.2.1 Current Implementation

- **Tabular Q-learning:** State space limited to 23,328 states; won't scale to large cities
- **Simulated environment:** Real-world deployment requires sensor integration and safety validation
- **Camera simulation:** YOLO detection is simulated; actual vision systems need hardware and calibration
- **Single objective:** Primary metric is waiting time; emissions and fuel consumption not modeled

5.2.2 Assumptions

- Perfect state observation (no sensor noise)
- Deterministic vehicle behavior (SUMO's Krauss model)
- No weather or road condition effects
- Fixed network topology (no construction or lane closures)

5.3 Lessons Learned

5.3.1 Technical Insights

1. **State representation matters:** Adding phase bits to the state eliminated aliasing and stabilized learning
2. **Reward shaping is critical:** We iterated through 5+ reward designs before achieving convergence
3. **Pretraining accelerates learning:** Imitation from Smart controller reduced training episodes by 40%

5.3.2 Practical Considerations

1. SUMO version compatibility is fragile (API breaks between releases)
2. Atomic file writes essential to prevent Q-table corruption during training
3. Dashboard visualization crucial for debugging RL agent behavior

6 Future Work

6.1 Short-Term Extensions

1. **Real Casablanca network:** Import OpenStreetMap data for intersections at Bachkou, Zerktouni, and 2 Mars
2. **Deep Q-Networks (DQN):** Replace tabular Q-learning with neural network approximation for scalability
3. **Multi-agent RL:** Decentralized control where each intersection is an independent agent with communication
4. **Emissions modeling:** Add CO₂ and fuel consumption to reward function

6.2 Long-Term Vision

1. **Real-world pilot:** Deploy at one intersection with safety override system
2. **Edge computing:** Run RL inference on local hardware (Raspberry Pi, NVIDIA Jetson)
3. **Cloud integration:** Central coordination server for city-wide optimization
4. **Public dashboard:** Real-time traffic visualization for city officials and citizens

7 Conclusion

This project successfully demonstrates that **reinforcement learning can outperform traditional traffic control methods** in realistic urban simulations. Our Q-learning agent achieved a **12.7% reduction in average waiting time** compared to fixed-time signals, while rule-based approaches (Smart and Vision) failed to improve upon the baseline.

Key contributions include:

1. Implementation of four distinct control strategies for comparative analysis
2. Design of a multi-intersection RL framework with joint action space
3. Demonstration of pretraining effectiveness for accelerated learning
4. Quantitative evaluation showing statistical significance of RL improvements

The results validate the feasibility of RL-based traffic control and provide a foundation for future real-world deployment in Moroccan cities. While current limitations (tabular Q-learning, simulation only) prevent immediate deployment, the project establishes clear pathways toward scalable, adaptive traffic management systems.

As urban populations continue to grow, intelligent traffic control will become increasingly critical. This work contributes to the broader vision of **Smart Cities** where AI-driven systems optimize urban infrastructure in real time, improving quality of life for millions of citizens.

References

References

- [1] Webster, F. V. (1958). *Traffic signal settings*. Road Research Technical Paper No. 39, Road Research Laboratory, London.
- [2] Hunt, P. B., Robertson, D. I., Bretherton, R. D., & Royle, M. C. (1981). *The SCOOT on-line traffic signal optimisation technique*. Traffic Engineering & Control, 23(4), 190-192.
- [3] Lowrie, P. R. (1990). *SCATS: Sydney co-ordinated adaptive traffic system—A traffic responsive method of controlling urban traffic*. Roads and Traffic Authority, New South Wales.
- [4] Zhang, H., Feng, S., Liu, C., Ding, Y., Zhu, Y., Zhou, Z., ... & Li, Y. (2019). *CityFlow: A multi-agent reinforcement learning environment for large scale city traffic scenario*. In The World Wide Web Conference (pp. 3620-3624).
- [5] Abdulhai, B., Pringle, R., & Karakoulas, G. J. (2003). *Reinforcement learning for true adaptive traffic signal control*. Journal of Transportation Engineering, 129(3), 278-285.
- [6] Wei, H., Zheng, G., Yao, H., & Li, Z. (2018). *IntelliLight: A reinforcement learning approach for intelligent traffic light control*. In Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining (pp. 2496-2505).
- [7] Chu, T., Wang, J., Codecà, L., & Li, Z. (2020). *Multi-agent deep reinforcement learning for large-scale traffic signal control*. IEEE Transactions on Intelligent Transportation Systems, 21(3), 1086-1095.
- [8] Lopez, P. A., Behrisch, M., Bieker-Walz, L., Erdmann, J., Flötteröd, Y. P., Hilbrich, R., ... & Wießner, E. (2018). *Microscopic traffic simulation using SUMO*. In 2018 21st International Conference on Intelligent Transportation Systems (ITSC) (pp. 2575-2582). IEEE.